

02_ForestProject-Modeling&Presentation

August 19, 2020

1 Prediciton (Classification) the Types of Trees.

1.0.1 This Jupyter Notebook contains;

- Classification Models for Predicting the Types of Trees,
- Visualization of the Result.

1.0.2 What I Plan to implement in terms of ML models :

- SVM (I will use `LinearSVC` model from `sklearn.svm` module),
- XGBoost (I will use `XGBClassifier` model from `xgboost` module)
- Additionally :
 - Decision Tree (I will use `DecisionTreeClassifier` model from `sklearn.tree` module)
 - KNN (I will use `KNeighborsClassifier` model from `sklearn.neighbors` module)
 - LGBM (I will use `LGBMClassifier` model from `lightgbm` module).

1.0.3 I will use `yellowbrick`, `seaborn` and `matplotlib` modules to visualize the model results.

1.0.4 Importing `covtype1.csv` dataset for modelling and required libraries.

```
[58]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sqlalchemy import create_engine
import warnings
from IPython.core.pylabtools import figsize
from scipy.stats import zscore
from scipy import stats
from numpy import percentile
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
```

```

from statsmodels.formula.api import ols
from scipy.stats import zscore
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC, LinearSVC
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
import seaborn as sns
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from sklearn.model_selection import TimeSeriesSplit
from yellowbrick.classifier import ClassificationReport
from yellowbrick.datasets import load_occupancy
from sklearn.metrics import f1_score
font_title = {'family': 'times new roman', 'color': 'darkred',
              'weight': 'bold', 'size': 14}

warnings.filterwarnings('ignore')
sns.set_style("whitegrid")

plt.rcParams['figure.dpi'] = 100

```

1.0.5 Reading dataset with pandas

```
[135]: df = pd.read_csv("covtype1.csv")
```

```
[8]: df.head()
```

```

[8]:   Elevation  Aspect  Slope  Horizontal_Distance_To_Roadways  Hillshade_9am  \
0      2596      51      3                                510            221
1      2590      56      2                                390            220
2      2804     139      9                                3180            234
3      2785     155     18                                3090            238
4      2595      45      2                                391             220

      Hillshade_Noon  Horizontal_Distance_To_Fire_Points  Wilderness_Area1  \
0                232                                6279                1
1                235                                6225                1
2                238                                6121                1
3                238                                6211                1
4                234                                6172                1

      Wilderness_Area2  Wilderness_Area3  ...  Soil_Type33  Soil_Type34  \
0                   0                  0  ...           0           0
1                   0                  0  ...           0           0

```

2	0	0	...	0	0
3	0	0	...	0	0
4	0	0	...	0	0

	Soil_Type35	Soil_Type38	Soil_Type39	Soil_Type40	Cover_Type \
0	0	0	0	0	5
1	0	0	0	0	5
2	0	0	0	0	2
3	0	0	0	0	2
4	0	0	0	0	5

	Square_Hypo_Distance	Average_Dist_Road_Hydro	Average_Elevation_Hydro
0	66564	384	1298
1	44980	301	1292
2	76049	1724	1434
3	72488	1666	1451
4	23410	272	1297

[5 rows x 46 columns]

1.0.6 Modeling

```
[23]: X = df.drop("Cover_Type", axis = 1)
```

```
[24]: y = df["Cover_Type"]
```

- Splitting the data set into two pieces : Test split - Train split

```
[25]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳ random_state=101)
```

1.0.7 XGBoost Classifier

```
[26]: xgb_classifier = XGBClassifier()
xgb_classifier.fit(X_train , y_train)
```

```
[26]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
importance_type='gain', interaction_constraints='',
learning_rate=0.300000012, max_delta_step=0, max_depth=6,
min_child_weight=1, missing=nan, monotone_constraints='()',
n_estimators=100, n_jobs=0, num_parallel_tree=1,
objective='multi:softprob', random_state=0, reg_alpha=0,
reg_lambda=1, scale_pos_weight=None, subsample=1,
tree_method='exact', validate_parameters=1, verbosity=None)
```

```
[27]: y_predicted = xgb_classifier.predict(X_test)
```

```
[31]: y_predicted
```

```
[31]: array([7, 3, 3, ..., 1, 3, 2], dtype=int64)
```

```
[44]: xgb_accuracy = accuracy_score(y_test, y_predicted)
```

```
[45]: # Very good!
```

```
xgb_accuracy
```

```
[45]: 0.8714599602298091
```

Now, I would like to show three plots for the results. - Class Prediction Error Bar Plot -
Confusion Matrix - Classification Report


```
[42]: from sklearn.datasets import make_classification
      from sklearn.model_selection import train_test_split
      from sklearn.ensemble import RandomForestClassifier
      from yellowbrick.classifier import ClassPredictionError

      visualizer = ClassPredictionError(xgb_classifier)

      # Fit the training data to the visualizer
      visualizer.fit(X_train, y_train)

      # Evaluate the model on the test data
      visualizer.score(X_test, y_test)

      # Draw visualization
      visualizer.show()
```



[42]: <matplotlib.axes._subplots.AxesSubplot at 0x2030637a048>


```
[41]: from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split as tts
from sklearn.linear_model import LogisticRegression
from yellowbrick.classifier import ConfusionMatrix

# The ConfusionMatrix visualizer takes a model
cm = ConfusionMatrix(xgb_classifier)

# Fit fits the passed model. This is unnecessary if you pass the visualizer a
#   ↳ pre-fitted model
cm.fit(X_train, y_train)

# To create the ConfusionMatrix, we need some test data. Score runs predict()
#   ↳ on the data
# and then creates the confusion_matrix from scikit-learn.
cm.score(X_test, y_test)

cm.show()
```




[41]: <matplotlib.axes._subplots.AxesSubplot at 0x203062dc390>

```
[43]: from sklearn.model_selection import TimeSeriesSplit
      from sklearn.naive_bayes import GaussianNB

      from yellowbrick.classifier import ClassificationReport
      from yellowbrick.datasets import load_occupancy

      visualizer = ClassificationReport(xgb_classifier, support=True)

      visualizer.fit(X_train, y_train)           # Fit the visualizer and the model
      visualizer.score(X_test, y_test)          # Evaluate the model on the test data
      visualizer.show()
```



[43]: <matplotlib.axes._subplots.AxesSubplot at 0x203065f2668>

- There is a little bit low **recall** score of the class 5. I concluded that it is caused by the values in the data set. %59.8 of the class 5 is predicted true. Although the **recall** score level of class 5 is a little bit low, **precision** and **f1** score is quite well.

1.0.8 LinearSVC

```
[111]: modelSVM = LinearSVC()
```

```
[112]: modelSVM.fit(X_train , y_train)
```

```
[112]: LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
verbose=0)
```

```
[114]: pred = modelSVM.predict(X_test)
```

```
[115]: SVM_accuracy = accuracy_score(pred, y_test)
```

```
[116]: SVM_accuracy
```

```
[116]: 0.23831968044709578
```

- It seems the model is failed. I decided to drop additional columns ('Average_Dist_Road_Hydro', 'Average_Elevation_Hydro') and re-split the dataset.

```
[68]: df['Square_Hypo_Distance'] = np.sqrt(df['Square_Hypo_Distance'])
df1 = df.drop(['Average_Dist_Road_Hydro', 'Average_Elevation_Hydro'], axis = 1)
X1 = df1.drop("Cover_Type", axis = 1)
y1 = df1["Cover_Type"]
```

- Splitting the data set into two pieces : Test split - Train split

```
[96]: X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1, test_size=0.3,
↳ random_state=101)
```

```
[118]: modelSVM1 = LinearSVC()
```

```
[119]: modelSVM1.fit(X1_train , y1_train)
```

```
[119]: LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
verbose=0)
```

```
[120]: pred1 = modelSVM1.predict(X1_test)
```

```
[121]: SVM_accuracy1 = accuracy_score(pred1, y1_test)
```

```
[122]: SVM_accuracy1
```

```
[122]: 0.5566145480054032
```

- I've doubled the **accuracy score** but it's still insufficient. SVM Classifier ;
 - SVM has been widely used in finance. For example, predicting stock price via SVM has been a acknowledged application in the industry.
 - In classification of text and handwritten objects, SVM performs well.
 - It may not be very successful in datasets with more than 100,000 data.

Now, I would like to show three plots for the results. - Class Prediction Error Bar Plot - Confusion Matrix - Classification Report

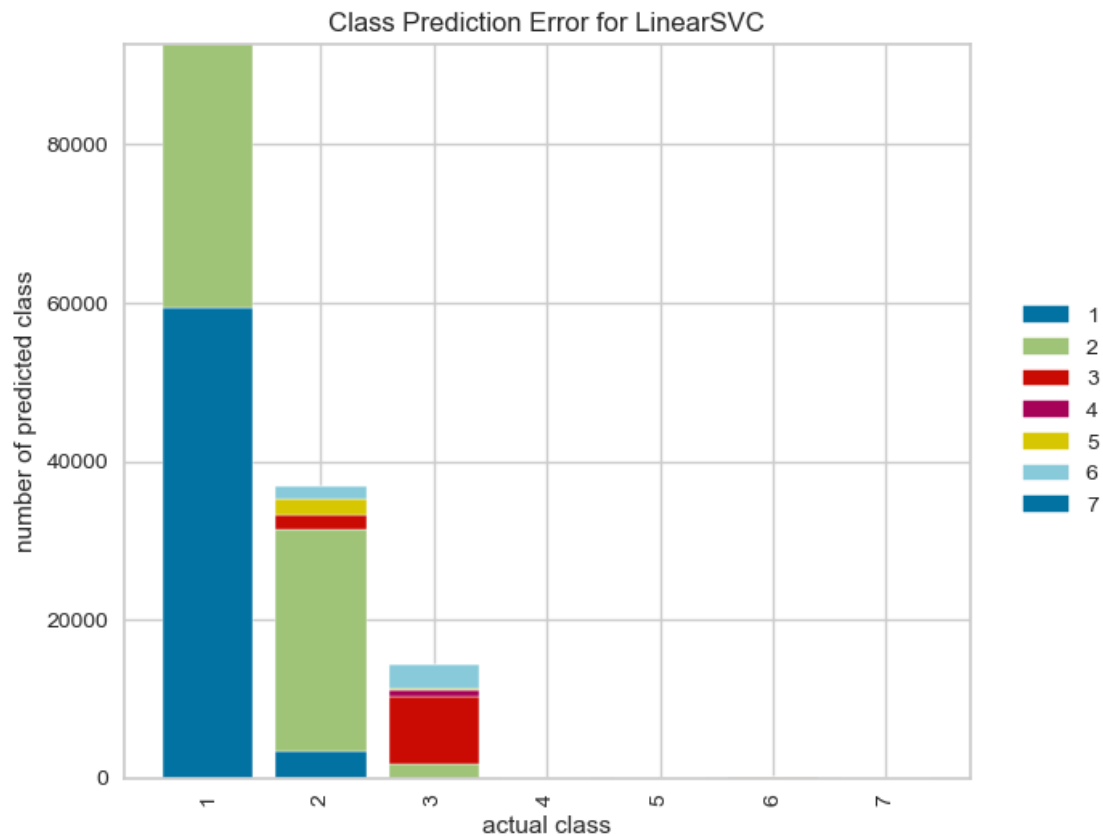
```
[123]: visualizer = ClassPredictionError(modelSVM1)

# Fit the training data to the visualizer
visualizer.fit(X1_train, y1_train)
```



```
# Evaluate the model on the test data
visualizer.score(X1_test, y1_test)

# Draw visualization
visualizer.show()
```




[123]: <matplotlib.axes._subplots.AxesSubplot at 0x2030b33fe10>

```
# The ConfusionMatrix visualizer takes a model
cm = ConfusionMatrix(modelSVM1)

# Fit fits the passed model. This is unnecessary if you pass the visualizer a
↳ pre-fitted model
cm.fit(X1_train, y1_train)

# To create the ConfusionMatrix, we need some test data. Score runs predict()
↳ on the data
# and then creates the confusion_matrix from scikit-learn.
cm.score(X1_test, y1_test)
```


```
cm.show()
```



```
[125]: <matplotlib.axes._subplots.AxesSubplot at 0x2030c4c5198>
```

```
[126]: visualizer = ClassificationReport(modelSVM1, support=True)

visualizer.fit(X1_train, y1_train)      # Fit the visualizer and the model
visualizer.score(X1_test, y1_test)      # Evaluate the model on the test data
visualizer.show()
```



[126]: <matplotlib.axes._subplots.AxesSubplot at 0x2030c5a59e8>

- We can see from the plots that more than half of the classes are predicted correctly. But the model predicted 5 classes although there are 7. Class 5 and 7 could not be detected.

1.0.9 DecisionTreeClassifier

```
[47]: modelTree = DecisionTreeClassifier()
```

```
[48]: modelTree.fit(X_train , y_train)
```

```
[48]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                             max_depth=None, max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort='deprecated',
                             random_state=None, splitter='best')
```

```
[49]: pred = modelTree.predict(X_test)
```

```
[50]: tree_accuracy = accuracy_score(pred, y_test)
```

```
[51]: # Quite well!
```

```
tree_accuracy
```

```
[51]: 0.9354517047266234
```

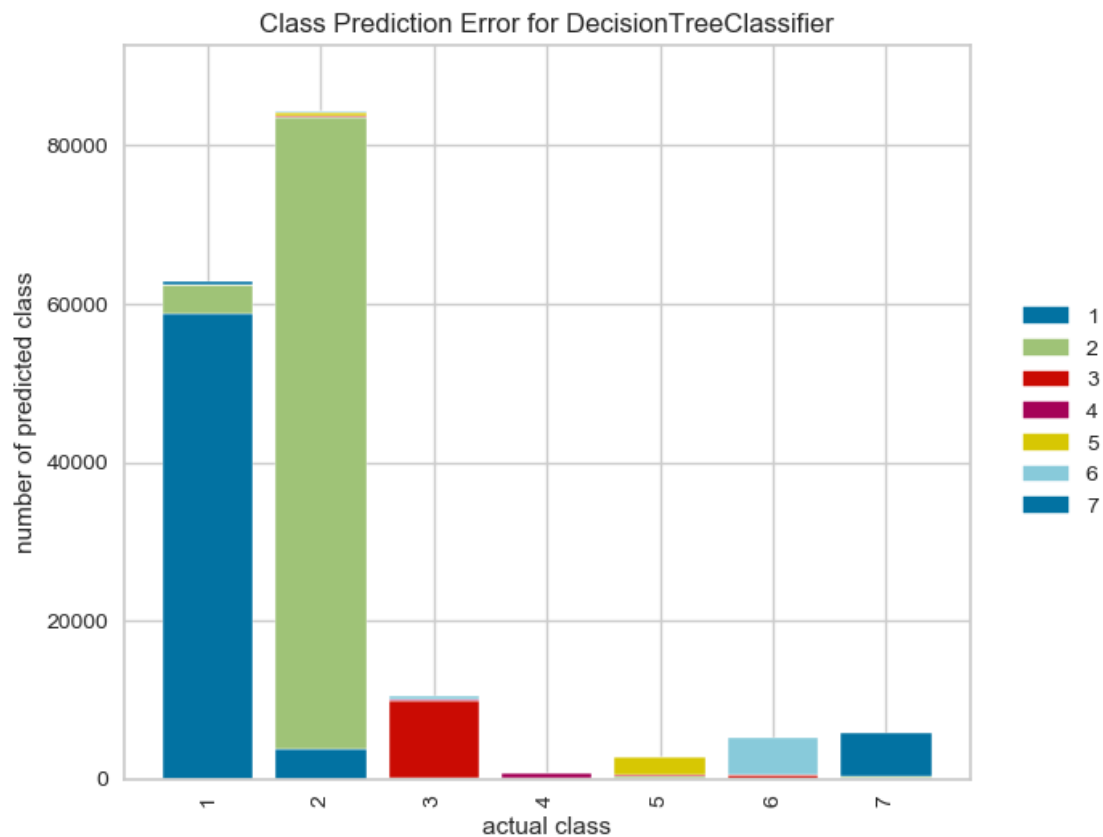
Now, I would like to show three plots for the results. - Class Prediction Error Bar Plot -
Confusion Matrix - Classification Report

```
[54]: visualizer = ClassPredictionError(modelTree)
```

```
# Fit the training data to the visualizer  
visualizer.fit(X_train, y_train)
```

```
# Evaluate the model on the test data  
visualizer.score(X_test, y_test)
```

```
# Draw visualization  
visualizer.show()
```



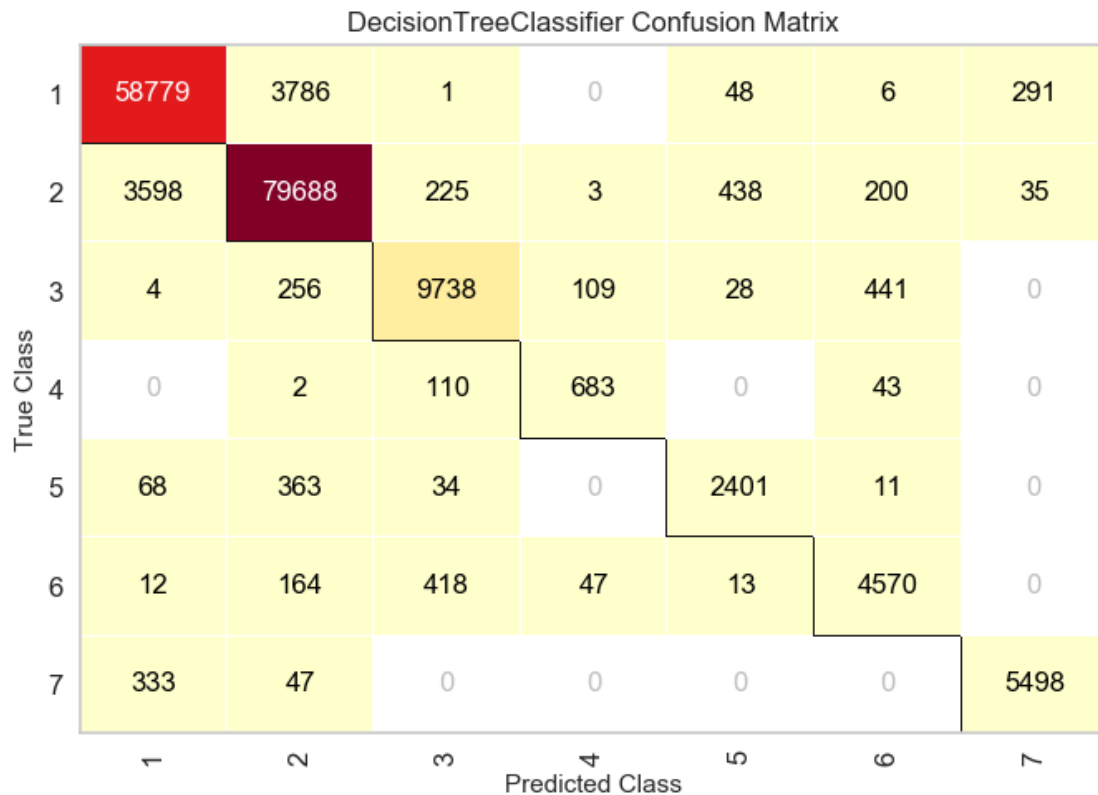
[54]: <matplotlib.axes._subplots.AxesSubplot at 0x203097713c8>

```
[53]: # The ConfusionMatrix visualizer takes a model
cm = ConfusionMatrix(modelTree)

# Fit fits the passed model. This is unnecessary if you pass the visualizer a
↳ pre-fitted model
cm.fit(X_train, y_train)

# To create the ConfusionMatrix, we need some test data. Score runs predict()
↳ on the data
# and then creates the confusion_matrix from scikit-learn.
cm.score(X_test, y_test)


cm.show()
```



[53]: <matplotlib.axes._subplots.AxesSubplot at 0x20309c15a20>

```
[52]: visualizer = ClassificationReport(modelTree, support=True)
```

```
visualizer.fit(X_train, y_train)      # Fit the visualizer and the model
visualizer.score(X_test, y_test)     # Evaluate the model on the test data
visualizer.show()
```



[52]: <matplotlib.axes._subplots.AxesSubplot at 0x20309c15940>

1.0.10 KNeighborsClassifier

- Deciding the number of neighbors

```
[55]: neighbors = np.arange(1, 7)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))

for i,k in enumerate(neighbors):
    #Setup a knn classifier with k neighbors
    knn = KNeighborsClassifier(n_neighbors = k)

    #Fit the model
    knn.fit(X_train, y_train)
```

```

#Compute accuracy on the training set
train_accuracy[i] = knn.score(X_train, y_train)

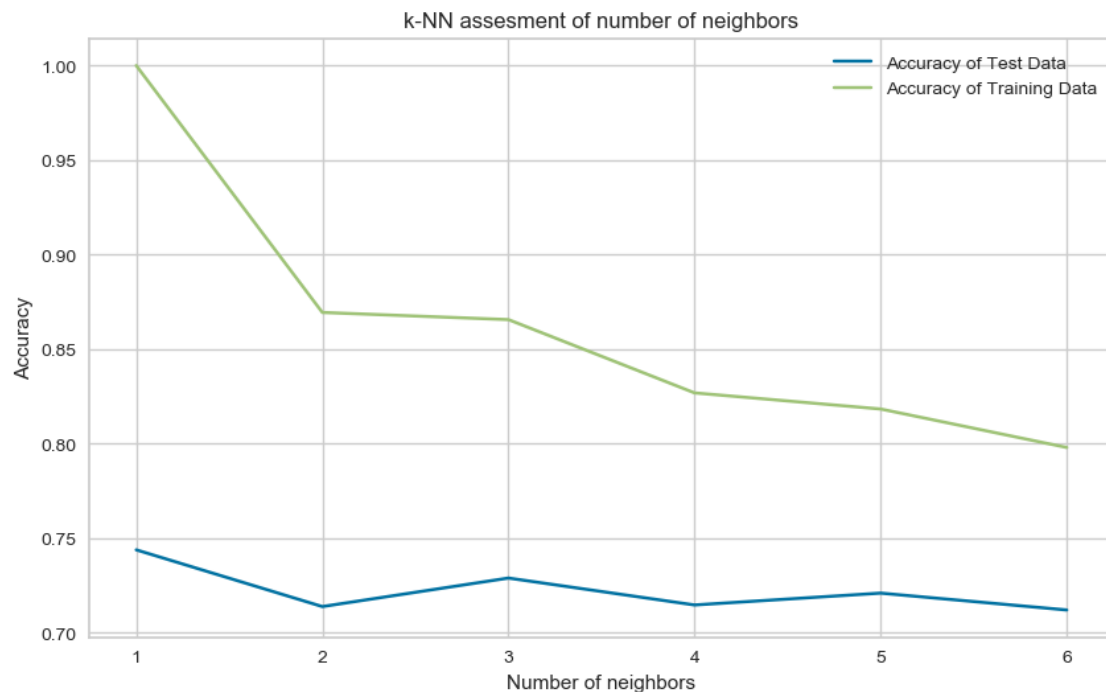
#Compute accuracy on the test set
test_accuracy[i] = knn.score(X_test, y_test)

```

```

[56]: plt.figure(figsize=(10,6))
plt.title('k-NN assesment of number of neighbors')
plt.plot(neighbors, test_accuracy, label='Accuracy of Test Data')
plt.plot(neighbors, train_accuracy, label='Accuracy of Training Data')
plt.legend()
plt.xlabel('Number of neighbors')
plt.ylabel('Accuracy')
plt.show()

```



- The graph lines stabilize around 5. So, Let's try 5 neighbors. Let's prepare the train and test data for KNN Model.

```

[101]: knn5 = KNeighborsClassifier(n_neighbors = 5)

```

```

[102]: knn5.fit(X1_train,y1_train)

```

```

[102]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=None, n_neighbors=5, p=2,

```

```
weights='uniform')
```

```
[103]: knn_accuracy = knn5.score(X1_test,y1_test)
```

```
[104]: # Excellent!
```

```
knn_accuracy
```

```
[104]: 0.962745882393864
```

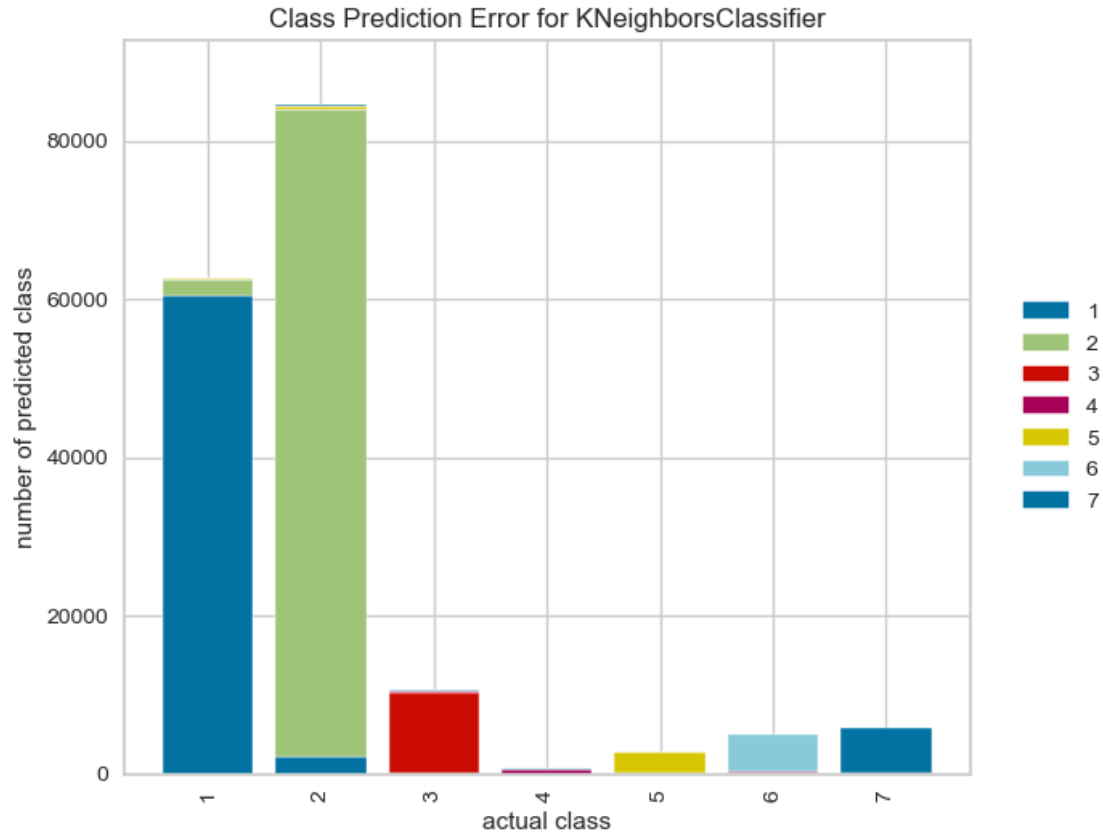
Now, I would like to show three plots for the results. - Class Prediction Error Bar Plot -
Confusion Matrix - Classification Report

```
[108]: visualizer = ClassPredictionError(knn5)

# Fit the training data to the visualizer
visualizer.fit(X1_train, y1_train)

# Evaluate the model on the test data
visualizer.score(X1_test, y1_test)

# Draw visualization
visualizer.show()
```




[108]: <matplotlib.axes._subplots.AxesSubplot at 0x20306be8358>

```
[109]: # The ConfusionMatrix visualizer takes a model
cm = ConfusionMatrix(knn5)

# Fit fits the passed model. This is unnecessary if you pass the visualizer a
# pre-fitted model
cm.fit(X1_train, y1_train)

# To create the ConfusionMatrix, we need some test data. Score runs predict()
# on the data
# and then creates the confusion_matrix from scikit-learn.
cm.score(X1_test, y1_test)


cm.show()
```



[109]: <matplotlib.axes._subplots.AxesSubplot at 0x203000f5358>

```
[110]: visualizer = ClassificationReport(knn5, support=True)

visualizer.fit(X1_train, y1_train)           # Fit the visualizer and the model
visualizer.score(X1_test, y1_test)          # Evaluate the model on the test data
visualizer.show()
```



```
[110]: <matplotlib.axes._subplots.AxesSubplot at 0x2030708e400>
```

1.0.11 LGBMClassifier

```
[77]: from lightgbm import LGBMClassifier

[80]: lgbm_classifier = LGBMClassifier()

[81]: lgbm_classifier.fit(X_train, y_train)
      y_predicted = lgbm_classifier.predict(X_test)

[127]: lgbm_accuracy = accuracy_score(y_test, y_predicted)
```

```
[128]: # Very good..
```

```
lgbm_accuracy
```

```
[128]: 0.8332782579960694
```

Now, I would like to show three plots for the results. - Class Prediction Error Bar Plot -
Confusion Matrix - Classification Report

```
[130]: visualizer = ClassPredictionError(lgbm_classifier)
```

```
# Fit the training data to the visualizer
```


```
visualizer.fit(X_train, y_train)
```

```
# Evaluate the model on the test data
```

```
visualizer.score(X_test, y_test)
```

```
# Draw visualization
```

```
visualizer.show()
```




```
[130]: <matplotlib.axes._subplots.AxesSubplot at 0x2030c5a5b38>
```

```
[131]: # The ConfusionMatrix visualizer takes a model
cm = ConfusionMatrix(lgbm_classifier)

# Fit fits the passed model. This is unnecessary if you pass the visualizer a
# pre-fitted model
cm.fit(X_train, y_train)

# To create the ConfusionMatrix, we need some test data. Score runs predict()
# on the data
# and then creates the confusion_matrix from scikit-learn.
cm.score(X_test, y_test)


cm.show()
```



```
[131]: <matplotlib.axes._subplots.AxesSubplot at 0x203096c2278>
```

```
[132]: visualizer = ClassificationReport(lgbm_classifier, support=True)

visualizer.fit(X_train, y_train)      # Fit the visualizer and the model
visualizer.score(X_test, y_test)      # Evaluate the model on the test data
visualizer.show()
```



[132]: <matplotlib.axes._subplots.AxesSubplot at 0x2032c4635f8>

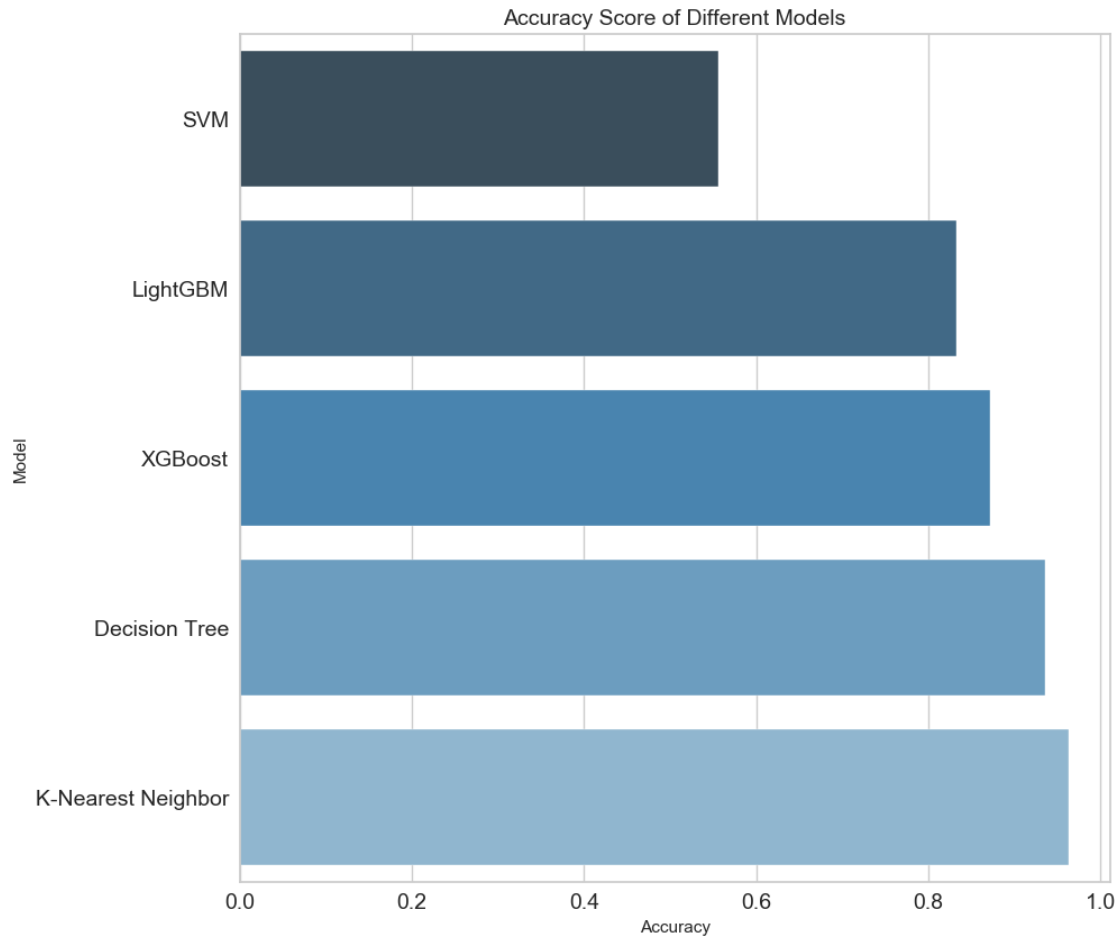
```
[134]: compare = pd.DataFrame({"Model": ["K-Nearest Neighbor", "LightGBM", "XGBoost",
    ↪ "Decision Tree", "SVM"],
    "Accuracy": [knn_accuracy, lgbm_accuracy, xgb_accuracy,
    ↪ tree_accuracy, SVM_accuracy1]})

compare = compare.sort_values(by="Accuracy", ascending=True)

plt.figure(figsize=(10,10))
ax = sns.barplot(x="Accuracy", y="Model", data=compare, palette="Blues_d")

plt.yticks(size = 14)
plt.xticks(size = 14)
plt.title("Accuracy Score of Different Models", size=14)
```

[134]: Text(0.5, 1.0, 'Accuracy Score of Different Models')



1.0.12 Overall Results :

- The overall evaluation in terms of **accuracy scores** is displayed above.
- The model with the least performance is SVM.
 - SVM has been widely used in finance. For example, predicting stock price via SVM has been a acknowledged application in the industry.
 - In classification of text and handwritten objects, SVM performs well.
 - It may not be very successful in datasets with more than 100,000 data (We have half a million).
 - It also doesn't perform well against unstable data (**Cover_Type**).
- **XGBoost** has performed well enough. As in **linear regression** models, it has proven its success in **clasification** once again.
 - There is a little bit low **recall** score of the class 5.
 - I concluded that it is caused by the values in the data set. %59.8 of the class 5 is predicted true.
 - Although the **recall** score level of class 5 is a little bit low, **precision** and **f1** score is quite well.

- All the other Models (K-Nearest Neighbor, LightGBM, Decision Tree) are quite well. They all have performed well and given high accuracy scores.
- Multi Class Classification is KNN's job...

Best Regards..